

# Gaussian Process Regression using GPyTorch

Lorena Magaña Zertuche

(Dated: April 7, 2022)

GstLAL, one of LIGO’s detection pipelines, relies on the use of matched-filtering. This allows us to get an estimate of the initial system parameters, such as masses, and spins. However, there is a significant discrepancy between the true values of the parameters and the recovered values. By using machine learning methods, such as Neural Networks and Gaussian Process Regression, one can decrease the errors between true and recovered values. A better parameter recovery means that we can better understand whether or not the remnant of a binary system will produce an electromagnetically bright event. This, in turn, will improve the reliability of the events sent out to astronomers. In this work, we show the results of using Gaussian Process Regression on the parameters of simulated (or injected) signals.

## I. INTRODUCTION

These document will outline the use of Gaussian Process Regression using GPyTorch and keep results as a means to keep this process as transparent as possible. Scikit-learn was explored earlier during the IPAM long program but its lack of flexibility when it comes to incremental learning and fine-tuning has made us move away from its further usage. Additionally, we looked into the GPR implementation of TensorFlow since it gives us the desired flexibility. However, both of these methods lack the speed of GPyTorch. For example, training and testing on the GstLAL fake data took approximately 4 hours for both Scikit-learn and TensorFlow. However, the same exact run took less than 4 minutes with GPyTorch. This is mainly due to the Lanczos Variance Estimates, or LOVE, method for fast variances and sampling introduced in (<https://arxiv.org/abs/1803.06058>).

## II. CONDITIONING THE DATASET

Thus far, GPR has been tested using a dataset generated from GstLAL early warning triggers. This low-latency pipeline uses matched-filtering techniques for the detection of gravitational-wave signals from compact binaries. The template bank of simulated binary neutron stars is generated using the TaylorF2 waveform approximant. The waveforms are for non-spinning systems with masses  $m_1 > 0.95M_\odot$  and  $m_2 < 2.4M_\odot$ .

Once the data is read, we prepare it for regression in the following way:

- (i) Shuffle the data.
- (ii) Constrain the data to ensure positive values in mass predictions. This requires
  - (a) Mapping the data to be in the range from (0, 1). This alone does not ensure positive values it is an intermediary step.
  - (b) Mapping the data to be in the range from  $(-\infty, +\infty)$ .
- (iii) Shuffle the data.




FIG. 1. The  $m_1$ - $m_2$  parameter space of injections is shown in red, the recovered values from matched-filtering are in blue, and the predicted values from regression are in green.

- (iv) Standardize the data so the data has a zero mean and unit variance.

Once the algorithm gives us the predicted data, the reverse process is done, i.e., the data is scaled back from standardization, mapped back to (0, 1), and exponentiated.

## III. METHODOLOGY

The kernel used for training is the radial basis function (RBF) which is a squared exponential function given by

$$K(X_1, X_2) = e^{-\frac{\|X_1 - X_2\|^2}{2\sigma^2}}, \quad (1)$$

where  $X_1$  and  $X_2$  are input data points and  $\sigma$  is the variance, or lengthscale parameter. The predictions of the GstLAL dataset are independent of kernel choice. However, the RBF kernel, apart from its popularity due to its versatility, seems to perform slightly better than other kernels and without a loss of speed.

The Adams optimizer is used to find the optimal hyper-parameters with a learning rate of 0.1 and 50 iterations. This combination of learning rate and iterations allows us to decrease the marginal log likelihood, or loss, at a quick rate without loss of accuracy.

Both the **GstLAL** and **O2** datasets are divided as 58% training and 42% testing.

#### IV. GSTLAL RESULTS




FIG. 2. The left panel show the injected (red line), recovered (blue stars), and predicted (green circles) values for  $m_1$  while the right panel shows the values for  $m_2$ . Regression is significantly better at producing the true values.

After running GPR on the datasets, we are able to better predict the masses of the binary systems. Fig. 2 shows the advantages of using regression. The values of the injected masses are depicted by the red line while the recovered masses are shown by the blue stars. The predicted masses, shown in green, are much more closely aligned in parameters space to the true values. Although the results from GPR are not perfect, the comparison between the recovered and predicted values is clear. Moreover, as illustrated in Fig. 3 the error in the recovery of the primary mass increases as the injected  $m_1$  increases and the error in the secondary mass decreases as the injected  $m_2$  value increases. Nevertheless, the error using the predicted values stays around zero regardless of the values of the masses injected.




FIG. 3. The left/right panel show the relative errors in the recovered (blue) and predicted (green) values for  $m_1/m_2$ . The red line is a fit of these values, which shows a relationship between injected  $m_1/m_2$  values and the recovered and predicted values. GPR has no significant bias that depends on the recovered mass while the **GstLAL** data does.

A different way to visualize Fig. 3 is by creating a histogram showing how the recovered and predicted values vary from those of the injected values. As seen in Fig. 4 the error in the recovered masses is very spread out with errors as high as 50%. However, using regression allows us to narrow down the errors to no more than about 10%. Overall, we are able to decrease the mean of the error by as much as three orders of magnitudes.




FIG. 4. The histogram on the left show the relative error between the injected  $m_1$  values and both the recovered (gray) and predicted (blue) values. Similarly, the right panel shows the relative error between the injected  $m_2$  values and both the recovered (gray) and predicted (orange) values. The mean of each dataset is shown by circles on the horizontal axis with the values in boxes. GPR does a great job at decreasing the mean error of the masses.

#### A. Statistical Analysis

**Shapiro-Wilk Test.** To better understand the dataset and the predictions, we perform some statistical tests. First, we check for the Gaussianity of the data by performing a Shapiro-Wilk test, i.e., by comparing the p-values of the injected, recovered, and predicted values, we are able to determine whether the results are normally distributed. Small p-values correspond to distributions that are not normal. As shown in Fig. 5, the injected masses are relatively normal with p-values of  $10^{-5}$  and  $10^{-12}$  for  $m_1$  and  $m_2$ , respectively. Instead the p-values for the recovered masses are on the order of  $10^{-38}$  and  $10^{-25}$  which indicates that the distribution cannot be considered “as normal” as the injected distribution. Furthermore, we use the same test for the predicted values and notice that the p-values resemble those of the injected masses. This tells us that using regression not only decreases the errors in the masses but also takes us back to a distribution similar to the injected one. Therefore, we can say that the predicted values agree with the null hypothesis of normality.

As an extra step, I was interested to see how the p-values change with the amount of data included. This is because for more than 5,000 data points, the p-value test may not be accurate. There are a couple of references to add here, but I will come back to that soon. Meanwhile, Fig. 6 show how the p-value changes as one add more data. Although the injected and predicted p-values seem to diverge slowly with an increasing number of data points, the values are




FIG. 5. The histograms show the mass distributions for the injected (left panel), recovered (middle panel), and predicted (right panel) masses. The p-values corresponding to the Shapiro-Wilk test for normality are printed in the plots. Note that the injected and predicted panels have similar p-values, showing that their distributions are similar in Gaussianity. On the other hand, the recovered p-values are extremely small. This means that they reject the null hypothesis of normality. The distributions are for  $N = 4,000$

still considered to be similar. Not surprisingly, the p-value of the recovered masses reject the null hypothesis that the data is normally distributed.

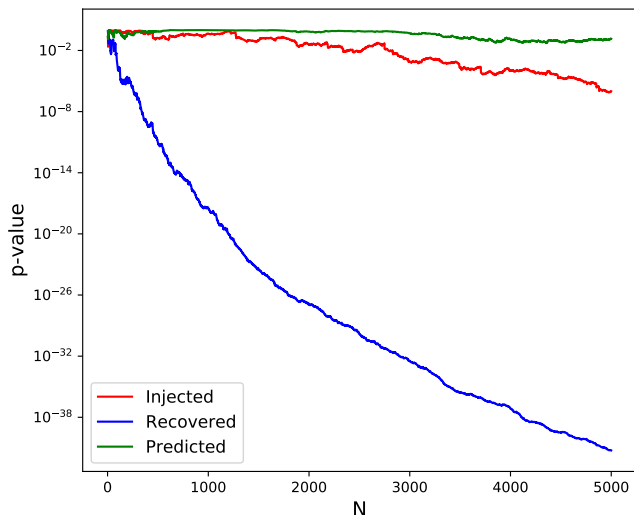


FIG. 6. The p-value as a function of number of data points is shown for the injected, recovered, and predicted data. Note that the p-values between the injected and predicted data sets are always closer together than those of the recovered data.

**Epps-Singleton Test.** The Epps-Singleton Test allows us to check if two sets of data have the same underlying probability distribution. A p-value of 1 means that two set of data have the same probability distribution while a p-value of 0 tells us that the two datasets are completely different in their distributions. When comparing the injected and recovered values, the p-value is  $2.5 \times 10^{-22}$ . On the other hand, comparing the injected and predicted values gives us a p-value of 0.32.

**Kolmogorov-Smirnov Test.** This last test compares the underlying continuous distributions of two indepen-

dent samples. What we find that that the p-value for the injected and recovered distributions is 0, i.e., the distribution of the injected values are not equal to the distribution of the recovered values. Instead for the predicted values, we get a p-value of 0.002 which tells us that although the continuous distributions are not the same, they are not completely different since there is some overlap. Fig. 7 shows the distribution of data as a function of mass. This visual shows us clearly that the recovered values for  $m_1$  tend to overestimate the mass and the recovered  $m_2$  values mostly underestimate the mass. This is corroborated by calculating the p-values. For example, the p-value between the injected and recovered  $m_1$  is 1 when the null hypothesis is that the underlying distribution of the injected parameters is less than the underlying distribution of the recovered parameters. In other words, the recovered masses overestimate the values. The converse gives us a p-value of 0. If instead we turn our attention to the right panel of Fig. 7, we can test for the distribution of the  $m_2$  values. The p-value between the injected and recovered  $m_2$  is 1 when the null hypothesis is that the underlying distribution of the injected parameters is greater than the underlying distribution of the recovered parameters. When comparing the predicted data, however, the p-values are small but not so small that any null hypothesis can be rejected. “Any” meaning that the continuous distribution of the predicted values is similar to the injected values.

## V. O2 DATA RESULTS

The same analysis will be done with the O2 data, which is currently running. Meanwhile, 1/7 of the data has been analysed and results do not look too good. However, better results are expected since these partial results represent a mass range that is 100 times larger than the previous test dataset. Stay tuned!




FIG. 7. The panels show the mass distributions for the injected, recovered, and predicted datasets. This hints at the biases in the recovered data.